

COMP 019 - Applications Development and Emerging Technologies

Emerging Technologies: MCP, A2A, Pgvector

SESSION 11: VECTOR DATABASES & PGVECTOR (EMERGING TECHNOLOGY)

ACTIVITY 11.1: INTRODUCTION TO VECTOR DATABASES

Topic: Understanding embeddings and vector similarity search

Description: Learn the fundamentals of vector databases for AI applications

INSTRUCTIONS:

Understand vector concepts:

- What are embeddings?
- Text → Vector representation
- Semantic similarity via vector distance
- Use cases: search, recommendations, RAG

Vector database options:

- Pgvector (PostgreSQL extension)
- Pinecone, Weaviate, Chroma (dedicated)
- Why Pgvector: Use existing PostgreSQL!

Embedding models:

- OpenAI text-embedding-ada-002
- Sentence Transformers (free, local)
- Google's embedding models

How semantic search works:

- Store text with embedding vectors
- Query: Convert question to vector
- Find similar vectors in database
- Return matching content

Deliverables: Submit vector database concepts document with diagrams

25 Points

ACTIVITY 11.2: SETTING UP PGVECTOR

Topic: Adding vector capabilities to PostgreSQL

Description: Install and configure Pgvector extension for vector search

INSTRUCTIONS:

Enable Pgvector:

Cloud PostgreSQL (Supabase/Neon):

- Usually pre-installed, just enable extension
- CREATE EXTENSION vector;

Local PostgreSQL:

- Install pgvector extension
- Enable in database

Create vector-enabled table:

- Add embedding column: vector(1536)
- Store: id, content, embedding
- Create index for fast search

Django integration:

- pip install pgvector
- Add VectorField to model
- Migration to add vector column

Generate embeddings:

- Use sentence-transformers or API
- pip install sentence-transformers
- Generate embedding for sample texts
- Store in database

Deliverables: Submit Pgvector setup code and Django model with vectors

35 Points

ACTIVITY 11.3: BUILDING RAG (RETRIEVAL-AUGMENTED GENERATION)

Topic: Creating AI-powered search with vector retrieval

Description: Build a RAG system combining vector search with AI generation

INSTRUCTIONS:

Understand RAG architecture:

- User asks question
- Convert question to embedding
- Search vector database for relevant content
- Send content + question to LLM
- LLM generates answer using retrieved content

Implement RAG system:

- Create embeddings for your content (student records, course materials)
- Store in Pgvector-enabled PostgreSQL
- Build search function
- Integrate with AI (Claude API or local model)

Build 'Smart Campus Assistant':

- Index: Course descriptions, policies, FAQs
- User asks: 'What are the requirements for the IT program?'
- System retrieves relevant documents
- AI generates helpful, accurate answer

Test with various queries

Compare: With RAG vs Without RAG

Deliverables: Submit RAG implementation code and comparison demo

40 Points

TOTAL POINTS: 100